More data manipulation

March 07, 2016

R packs a lot of functionality to make data manipulation easier, and it is impossible to even give an overview of the array of tools that are available in this little space. However, the majority of the things I end up doing fall in three categories: apply some function to a grouping of the data, merging separate datasets, or reshaping data. Moreover, I think they summarize nicely they way R works.

Split-apply-combine

Consider again our tobacco example and the problem that we had about calculating group means. We used a specific function, `ave` that did precisely that. But the operation is actually a lot more general. We can think about spliting data structures by gropus, performing operations in each group and then grouping things again. Base R provides functions to accomplish this kind of tasks, but I personally prefer the functions in the `plyr` package, because they provide a common, clear, and intuitive interface.

The usage is very similar across functions. The first argument is the input data (such as a `data.frame` or a `list`). Then, a variable that tells how the input should be split into groups, although it may not be needed –think about a `list`. Finally, a description of the operation to be applied to each group. The logic will become clearer with a couple of examples. Let's start by realizing that `plyr` is not shipped with R. After installing it, we can load it and also read in some data.

```
library(plyr)
tobacco <- read.csv("http://koaning.io/theme/data/cigarette.csv")
```

Let's calculate the average income per state (across years). The name of the function gives a lot information: the first letter `d` says that the input will be a `data.frame`; the second letter, that the output will be another data.frame. All functions in `plyr` use the same input-output structure:

```
group_means <- ddply(tobacco, ~ state, mutate, avincome=mean(income))
head(group_means)
```

```
##   state year  cpi      pop packpc   income  tax avgprs taxs avincome
## 1    AL 1985 1.08 3973000    116 46014968 32.5    102 33.3 64137227
## 2    AL 1986 1.10 3992000    117 48703940 32.5    108 33.4 64137227
## 3    AL 1987 1.14 4016000    116 51846312 32.5    114 33.5 64137227
## 4    AL 1988 1.18 4024000    115 55698852 32.5    120 33.5 64137227
## 5    AL 1989 1.24 4030000    109 60044480 32.5    133 33.7 64137227
## 6    AL 1990 1.31 4048508    112 64094948 32.5    143 33.8 64137227
```

See how the grouping is passed as a formula, and now we are giving a name `avincome` to the newly created variable, which is the mean of the `income` variable. The `mutate` argument indicates that we want to perserve the number of rows in the dataset. Compare the output with what happens when we use `summarize` which says that we only one one value per group:

```
group_means <- ddply(tobacco, ~ state, summarize, avincome=mean(income))
head(group_means)
```

```
##    state avincome
```

```
## 1    AL 6.41e+07
## 2    AR 3.50e+07
## 3    AZ 6.41e+07
## 4    CA 6.23e+08
## 5    CO 6.71e+07
## 6    CT 8.45e+07
```

Let's explore a bit more complex example that calculates a separate regression between `packpc` and `log(tax)` for each state and then pulls all the coefficients together.

```
lm_models <- dlply(tobacco, ~ state, function(x) lm(packpc ~ log(tax), data=x))
lm_coefs <- ldply(lm_models, coefficients)
head(lm_coefs)
```

```
##    state (Intercept) log(tax)
## 1    AL         304    -54.7
## 2    AR         282    -43.8
## 3    AZ         242    -43.1
## 4    CA         226    -40.2
## 5    CO         397    -84.1
## 6    CT         302    -52.6
```

Several things are worth noting. First, that we are first transforming a `data.frame` into a `list` (`dlply`), and then a `list` into a `data.frame` (`ldply`). Can you see why? Second, that the function in our `dlply` called is passed as a anonymous function which takes the data for each country as argument. This is a very common strategy in R: we need a temporary function for a one-off task and we don't even need a name for it. In this case, we just want a function that applies the same model to different datasets. Finally, that `ldply` only needs two arguments: because `lm_models` is already a list, we don't need to do any splitting.

Merging two datasets

From the moment a `data.frame` is an object, we can hold as many as our computer allows. And we can put them together and merge them by specifying the keys that they have in common. Consider a trivial example in which we want to take our original dataset and merge into it the vector of means by state that we calculated above.[1]

```
merged_tobacco <- merge(tobacco, group_means, by="state")
head(merged_tobacco)
```

```
##   state year  cpi      pop packpc   income  tax avgprs taxs avincome
## 1    AL 1985 1.08 3973000    116 46014968 32.5    102 33.3 64137227
## 2    AL 1992 1.40 4139269    107 72281824 36.5    176 38.1 64137227
## 3    AL 1991 1.36 4091025    107 67649568 34.5    162 35.9 64137227
## 4    AL 1987 1.14 4016000    116 51846312 32.5    114 33.5 64137227
## 5    AL 1986 1.10 3992000    117 48703940 32.5    108 33.4 64137227
## 6    AL 1989 1.24 4030000    109 60044480 32.5    133 33.7 64137227
```

The `merge` function takes some other arguments to specify different classes of joins, what to do with the units that don't match, or what to do with duplicated variables.

---

[1]Yes, we trying to accomplish the same thing we did in the `mutate` call.

Reshaping a dataset

We can change the structure and reshape our dataset so that rather than having state-year observations, each row represents data across years. Again, there is a very direct way to do it with base R, but I like the way the `reshape2` package works. Operations are split into two separate functions. The first one "melts" the dataset according to some indexing variables. The second one "casts" the data into a particular shape using a formula interface. The functions are very aptly named `melt` and `dcast`:

```r
library(reshape2)
molten_tobacco <- melt(tobacco[, c("state", "year", "tax")], id=c("state", "year"))
head(dcast(molten_tobacco, state ~ variable + year))
```

```
##    state tax_1985 tax_1986 tax_1987 tax_1988 tax_1989 tax_1990 tax_1991
## 1     AL     32.5     32.5     32.5     32.5     32.5     32.5     34.5
## 2     AR     37.0     37.0     37.0     37.0     37.0     37.0     39.0
## 3     AZ     31.0     31.0     31.0     31.0     31.0     31.0     35.2
## 4     CA     26.0     26.0     26.0     26.0     38.5     51.0     53.0
## 5     CO     31.0     31.0     36.0     36.0     36.0     36.0     38.0
## 6     CT     42.0     42.0     42.0     42.0     45.5     56.0     58.0
##    tax_1992 tax_1993 tax_1994 tax_1995
## 1      36.5     38.5     40.5     40.5
## 2      42.0     49.2     55.5     55.5
## 3      38.0     40.0     42.0     65.3
## 4      55.0     57.0     60.0     61.0
## 5      40.0     42.0     44.0     44.0
## 6      63.8     67.0     71.0     74.0
```

We don't always need to perform the two steps and if you look at the output of `molten_tobacco` you will realize that it is not doing much for us other than adding a factor that holds the variable *not* indicated in the `id` argument, and putting its values in a separate column.[2]

---

[2] We could have omitted the first step and run `dcast(tobacco[, c("state", "year", "tax")], state ~ year)`.